
Exploiting Convolutional Neural Network for Risk Prediction with Medical Feature Embedding

Zhengping Che

University of Southern California
zche@usc.edu

Yu Cheng

IBM T. J. Watson Research Center
chengyu@us.ibm.com

Zhaonan Sun

IBM T. J. Watson Research Center
zsun@us.ibm.com

Yan Liu

University of Southern California
yanliu.cs@usc.edu

1 Introduction

The widespread availability of electronic health records (EHRs) promises to usher in the era of personalized medicine. However, the problem of extracting useful clinical representations from longitudinal EHR data, sometimes called computational phenotyping, remains challenging, owing to the heterogeneous, longitudinally irregular, noisy and incomplete nature of such data.

In this paper, our focus is on the problems of high dimensionality and temporality. We explore deep neural network models with learned medical feature embedding to deal with these issues. Specifically, we use a multi-layer convolutional neural network (CNN) to parameterize the model and is thus able to capture complex non-linear longitudinal evolution of EHRs. Different from recent proposed deep learning approaches such as stacked auto-encoders [11] and recurrent neural network [5, 2], our model can effectively capture local/short temporal dependency in EHRs, which is beneficial for risk prediction. To account for high dimensionality, instead of using the raw EHR data as the input, we use the embedding medical features in the CNN model. Based on the medical context, each medical event is compressed into a given length vector with medical feature embedding. Similar to the word embedding [8], the event embedding presented in our model holds its natural medical concept. Our initial experiments produce promising results, and demonstrate the effectiveness of both the medical feature embedding and the proposed convolutional neural network in risk prediction on cohorts of congestive heart failure and diabetes patients, compared with several strong baselines.

2 Methodology

2.1 Medical Feature Embedding Learning

We extend word2vec model [8, 3] to learn medical feature embeddings from EHR data. Word2vec takes a corpus of text sequences and learns the vector representations, called embeddings, for each word. The words with similar contexts usually have close embeddings in the vector space. In healthcare domain, Word2vec has also been used for medical texts and claims [4]. Since our EHRs are also sequential data with thousands of different medical events (e.g., diagnosis, medications, lab test, etc.), it is natural to generalize the ideas in word2vec to learn low-dimension and meaningful embeddings for medical events given a large corpus of EHR data. We take all the records for one patient as a sequence of medical events. For each event in the sequence, the word2vec model takes other events within a local window of current event into consideration. We train word2vec with continuous bag-of-words (CBOW) method.




Figure 1: Convolutional neural network prediction model (with filters of size 2 and 3).

2.2 CNN Risk Prediction Model

Convolutional neural network (CNN) is one type of neural network models which is good at capturing local structure and has moderate model size. A variety of CNN models are shown to be effective in handling images [7], sequential data [6], and recently on medical lab tests [9]. In risk prediction tasks on EHR data of diagnoses and medications, the temporal structures and local dependencies across features also provide lots of useful information. In order to effectively learn these useful characteristics from EHR data and make better predictions, we use a **CNN with one-dimensional convolutional operation over the temporal dimension**. The input to our CNN model is the EHR records of patient p , which is represented as a temporal embedding matrix $\mathbf{X}_p \in \mathbb{R}^{T_p \times D}$, where T_p , which is usually different among the patients, is the number of medical events in patient p 's record, and D is the learned embedding dimension. Each row $\mathbf{x}_i \in \mathbb{R}^D, i = 1, \dots, T_p$ in \mathbf{X}_p is the embedding vector of i th event of that patient. Noticing that no spatial or temporal relationship inherently exists along the embedding dimension in the EHR matrix, we apply convolutional operation not over embedding dimension but only over the temporal dimension. Assume that the CNN model has K convolutional filters of size F , we will have K different vector outputs of length $T_p - F + 1$. A filter of size 2 captures pairwise temporal dependency, while a filter of larger size might capture longer temporal dependency. That said, using a combination of filters with different lengths benefits capturing dependencies in multiple levels and thus improves the prediction, which is validated in our experiments. After convolutional step, we apply a max-pooling operation along the temporal dimensions to keep the most important features across the time. This temporal pooling produces output vector of length K and moreover converts the inputs with different temporal lengths into a fixed length vector inherently. Finally a fully connected softmax layer is used to get the prediction. An illustration of the proposed CNN prediction model is shown in Figure 1. Our convolutional network model is similar to the basic model shown in [1] but with multiple scale filters. Furthermore, the input to our model is temporal embedding matrix of events instead of raw event matrix. We can also jointly fine-tune a task-specific embedding matrix together with training the CNN model, or use both fixed and fine-tuned embedding together as input matrix.

3 Experiments

3.1 Experimental Setup

3.1.1 Medical Feature Embedding

We applied medical feature embedding learning method described in Section 2.1 from a private EHR dataset. This dataset contains 218680 patients and 14969489 observations of 14690 unique medical events, which are ICD-9 diagnosis and medications. We only compute the embeddings for 8627 events that show at least 5 times in this dataset. We set the local window size to be 20 for each event, which means at most 40 events nearby are considered to be the neighbors of that event. We set the embedding dimension is 200 in all the following experiments, if not mentioned explicitly. Our medical feature embedding model is implemented based on python Gensim [10] package.

3.1.2 Risk Prediction

We apply the convolutional neural network and the learned embeddings to **predict two important diagnoses, diabetes and congestive heart failure**, as two binary classification tasks separately. We take the same private EHR dataset used in medical embedding learning, but only a small portion of patients in the dataset have the target diagnoses. In order to handle this class imbalance issue and compare all methods robustly, we respectively **extract subsets for these two diagnosis**. First, all patients with the target diagnosis are selected to build the case group. We take both the ICD-9 diagnoses and

Table 1: Prediction performance comparison of convolutional neural network and baselines.

Method	Input	Heart Failure		Diabetes	
		Accuracy	AUROC	Accuracy	AUROC
CNN	W2v	0.8630	0.9329	0.9844	0.9989
CNN	RAND	0.8337	0.8999	0.9815	0.9959
CNN	RAW	0.8511	0.9203	0.9748	0.9953
LR	BoFW	0.8094	0.8652	0.9066	0.9681
SVM	BoFW	0.7906	0.8397	0.9044	0.9687
RF	BoFW	0.8620	0.9311	0.9029	0.9683
LR	W2v	0.8625	0.9289	0.9266	0.9802
SVM	W2v	0.8476	0.9140	0.9148	0.9753
RF	W2v	0.8650	0.9269	0.8955	0.9654

medications as input features for heart failure prediction, but only diagnoses for diabetes since some medications may apparently indicate diabetes. If there are multiple events at the same time, we generate the input sequence by ordering them ascendingly by the event index. As patients usually have medical records in different lengths, it might impact the prediction performance and prevent us from learning robust model and comparing models consistently. We remove all patients with less than 50 medical event records, and keep the first up to 250 events before the time they get their first target diagnosis. For each patient in the case group, two other patients without the target diagnosis but with similar demographic information and EHR record length are selected to form a control group. Similarly we limit the length of each record sequence between 50 and 250. Finally we have 2248 and 4496 patients in case and control group for diabetes, and 3357 and 6714 patients in case and control group for heart failure. We split all the data into training, validation, and test subsets by 7:1:2. The CNN model we used in experiments has one convolutional layer with 100 filters of size 3, 4 and 5, followed by one pooling layer and one fully connected prediction layer. We take tanh function as activation in convolutional and fully connected layer since it performed better in our experiments. We train the network using AdaDelta [12] with default settings.

3.2 Experimental Results

In our experiments, we take classification accuracy (Accuracy), the area under receiver operating characteristic curve (AUROC), the area under precision-recall curve (AUPRC), and maximum F1 score (Max F1) to compare the performance of all methods in two risk prediction tasks.

3.2.1 Risk Prediction Comparison

First we want to show the performance of our proposed predict model. We compare logistic regression (LR), linear support vector machine (SVM) and random forest (RF) as baselines. We apply L2 regularizers in LR and SVM. We use early stopping for RF with at most 50 trees. We evaluate baselines with either raw inputs or learned embeddings. For raw input (BoFW), we take bag of words to keep only the frequencies of all events, thus the input vector is of length 8627. For word2vec embedding input (W2v), we take the embeddings of all events in that record. In order to make the patient records in different temporal lengths to be the same length. we take the columnwise aggregation values along the temporal dimension, including summation, maximum, and minimum of all embeddings in the record, which is equivalent to W2v-ALL shown in Section 3.2.2. In this way, the input length is related to only the embedding dimension and much smaller than raw input length. For CNN model, besides the learnt medical feature embeddings (W2v), we also evaluate the performance with raw feature index as input(RAW), and use random vector to initialize embedding and jointly trained embeddings together with CNN model (RAND). Table 1 shows the classification accuracy and AUROC on the two prediction tasks. The proposed CNN methods with learned embeddings is among the best methods in heart failure prediction task and significantly outperforms baselines in diabetes prediction task. The performance improvement mainly comes from the learned embeddings for heart failure task, but from CNN model structures for diabetes task. The learned embedding helps LR and SVM a lot, but RF barely benefits from the it, which is probably because RF selects discrete input feature at each tree node.

Table 2: Learned embeddings v.s. raw inputs on logistic regression prediction model.

	Heart Failure				Diabetes			
	Accuracy	AUROC	AUPRC	Max F1	Accuracy	AUROC	AUPRC	Max F1
BOFW	0.8094	0.8652	0.7590	0.7302	0.9066	0.9681	0.9386	0.8654
RAND-SUM	0.8005	0.8784	0.7818	0.7382	0.8488	0.9317	0.8869	0.8030
W2V-AVE	0.8471	0.9119	0.8133	0.7792	0.8977	0.9583	0.8848	0.8639
W2V-SUM	0.8407	0.9142	0.8389	0.7797	0.9266	0.9802	0.9638	0.8950
W2V-MAX	0.8422	0.9157	0.8267	0.7878	0.8695	0.9431	0.8654	0.8168
W2V-ALL	0.8625	0.9289	0.8668	0.8056	0.9162	0.9740	0.9530	0.8763

3.2.2 Feature Embedding Evaluation

Next, we compare the learned medical feature embeddings and raw representations by logistic regression prediction models. Comparing with raw feature index BOFW, we take several representations based on learned embeddings: W2V-AVE, W2V-SUM, W2V-MAX takes the average, summation, and maximum values of the columnwise embedding inputs, respectively. The input size is 200 for these three methods. We also take a combination of summation, minimum and maximum values, which is used in Section 3.2.1, as W2V-ALL. Additionally, we take the summation of random embeddings (RAND-SUM) as comparison. As shown in Table 2, our learned word2vec embeddings outperform the raw input and random embedding, especially for heart failure prediction.

3.2.3 Early Prediction

Making better prediction in advance can help doctor make timely decisions and diagnose diseases in their early stage. We also test our model in a simulated early stopping setting. For each patient in the case group, we only take the observations before a period of time from the time that the first target diagnosis is confirmed to the patient. We set the hold-off period to be 90/180 days in our experiments, to predict whether the patient will have the target diagnosis in 90/180 days later. Some of the patients do not have enough records before the first target diagnosis and thus are removed from the dataset. Thus the longer hold-off period we select, the fewer patients will be taken in the case group which makes the problem more difficult. We also remove some patients in control group to make sure the size of control group is always double of that of case group. The number of patients and early prediction performance is shown Table 3. As expected, the performance decreases when the duration goes up, but it’s still comparable with baselines shown in Table 1. For instance, our 180-day early prediction on diabetes achieves Accuracy of 0.9277, which is even higher than the that of best baseline with full observation, which accuracy of 0.9266 from LR+W2V-ALL.

4 Discussion

In this work, we proposed a general framework for risk prediction tasks by convolutional neural network and learned medical feature embedding on EHR data. We found that the learned embeddings of are better than raw representations in both baselines and our proposed CNN models. We also evaluated our framework in details and demonstrated its capability to capture complex temporal structures and make reasonable early predictions. In the future we plan to consider temporal gaps of events in embedding and prediction model, and find clinical interpretations of the learned embeddings and filters in the network.

Table 3: Early prediction results for CNN prediction model.

	Heart Failure			Diabetes		
	# of Case	Accuracy	AUROC	# of Case	Accuracy	AUROC
0 days	3357	0.8630	0.9329	2248	0.9844	0.9989
90 days	2573	0.8329	0.8889	1616	0.9835	0.9988
180 days	2105	0.8274	0.8995	1221	0.9277	0.9716

References

- [1] Y. Cheng, F. Wang, P. Zhang, and J. Hu. Risk prediction with electronic health records: A deep learning approach. *SDM*, 2016.
- [2] K. Cho, B. Van Merriënboer, D. Bahdanau, and Y. Bengio. On the properties of neural machine translation: Encoder-decoder approaches. *arXiv preprint arXiv:1409.1259*, 2014.
- [3] E. Choi, M. T. Bahadori, E. Searles, C. Coffey, M. Thompson, J. Bost, J. Tejedor-Sojo, and J. Sun. Multi-layer representation learning for medical concepts. In *Proceedings of the 22Nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, KDD '16, pages 1495–1504, New York, NY, USA, 2016. ACM.
- [4] Y. Choi, C. Y.-I. Chiu, and D. Sontag. Learning low-dimensional representations of medical concepts. *AMIA Summits on Translational Science Proceedings*, 2016:41, 2016.
- [5] S. Hochreiter and J. Schmidhuber. Long short-term memory. *Neural computation*, 9(8):1735–1780, 1997.
- [6] N. Kalchbrenner, E. Grefenstette, and P. Blunsom. A convolutional neural network for modelling sentences. *arXiv preprint arXiv:1404.2188*, 2014.
- [7] A. Krizhevsky, I. Sutskever, and G. E. Hinton. Imagenet classification with deep convolutional neural networks. In *Advances in neural information processing systems*, pages 1097–1105, 2012.
- [8] T. Mikolov and J. Dean. Distributed representations of words and phrases and their compositionality. *Advances in neural information processing systems*, 2013.
- [9] N. Razavian and D. Sontag. Temporal convolutional neural networks for diagnosis from lab tests. *arXiv preprint arXiv:1511.07938*, 2015.
- [10] R. Řehůřek and P. Sojka. Software Framework for Topic Modelling with Large Corpora. In *Proceedings of the LREC 2010 Workshop on New Challenges for NLP Frameworks*, pages 45–50, Valletta, Malta, May 2010. ELRA. <http://is.muni.cz/publication/884893/en>.
- [11] P. Vincent, H. Larochelle, I. Lajoie, Y. Bengio, and P.-A. Manzagol. Stacked denoising autoencoders: Learning useful representations in a deep network with a local denoising criterion. *Journal of Machine Learning Research*, 11(Dec):3371–3408, 2010.
- [12] M. D. Zeiler. Adadelta: an adaptive learning rate method. *arXiv preprint arXiv:1212.5701*, 2012.